Geography-Aware Province-Level Vietnamese Dialect Identification with Frozen Self-Supervised Speech Representations

Tho Dang
University of Information Technology, VNU-HCM
Ho Chi Minh City, Vietnam
Student ID: 24521715

***Abstract*—We study province-level Vietnamese dialect identification on ViMD, a 102-hour corpus of interviews broadcast by the television stations of all 63 provinces. Instead of fine-tuning, we keep self-supervised speech encoders frozen and train a light head over a learned weighted sum of their layers. A head over fused XLS-R and PhoWhisper representations reaches 63.9% accuracy as a three-seed ensemble (61.0% for individual seeds), and an ensemble of nine heads reaches 64.6% accuracy and 0.637 macro-F1 on the official test set; the best fine-tuned baseline reported with the dataset obtained 0.411 macro-F1 under a different protocol. We also study Geo-aware Label Smoothing (Geo-LS), a simple variant of distance-based soft labels that keeps a fixed mass on the true province and spreads the smoothing mass to geographically close provinces. On the fused model, Geo-LS lowers the mean distance between the true and predicted province by 8.2 km (95% CI 1.4 to 14.6) relative to standard label smoothing, with an accuracy difference of -0.3 points (95% CI -1.5 to 0.9). Control analyses indicate that the accuracy is not explained by spoken place names or by the transcript and long-term spectral cues that simple baselines exploit; because each label denotes the province of the broadcasting station, residual recording-source effects cannot be excluded. In the configurations we tried, full fine-tuning, LP-FT and temporal CNN or Transformer heads did not outperform the frozen-feature approach in this low-resource setting of about 240 training utterances per province.**

***Index Terms*—dialect identification, Vietnamese, self-supervised learning, XLS-R, Whisper, label smoothing, geographic distance**

I. Introduction

Vietnamese is usually described with three dialect regions, Northern, Central and Southern, but pronunciation also varies from province to province. Knowing a speaker's dialect helps speech recognition, speaker profiling and sociolinguistic research. The ViMD corpus [1] made the fine-grained task measurable by labelling speech from all 63 provinces. Its baselines fine-tune wav2vec 2.0 and Whisper models and reach at most 0.411 macro-F1.

The task has two properties that standard classification ignores. First, the classes have a geography: confusing Hà Nội with Hà Nam is a much smaller error than confusing it with Cà Mau. Second, a province label can be predicted from cues that are not dialect at all, such as place names spoken in the recording or the acoustic signature of the local broadcaster that recorded it. A useful study has to measure both.

This paper asks four research questions:

RQ1 How far can frozen self-supervised speech representations go on 63-way dialect identification?

RQ2 To what extent can the accuracy be explained by spoken content or recording channel rather than accent?

RQ3 Can a training objective that knows the map make errors geographically closer without a large loss in accuracy?

RQ4 Do fine-tuning or temporal architectures add anything over pooled frozen features at this data size?

Our contributions are:

- A frozen-encoder pipeline with a learned layer-weighted head that reaches 64.6% accuracy and 0.637 macro-F1 on the ViMD test set, about 0.23 macro-F1 above the published fine-tuned baselines, which were obtained under a different protocol (Section VI-D).
- A set of confound controls, including text-only, spectrum-only and place-name-free evaluations. They show that explicit place names and the cues captured by these baselines explain little of the accuracy; they cannot fully separate accent from recording source.
- Geo-aware Label Smoothing (Geo-LS), a simple adaptation of distance-based soft labels [2], [3] that fixes the target mass on the true class. Relative to standard label smoothing, it reduced the mean geographic error by 8 to 19 km in all four settings tested.
- Negative results with diagnostics for full fine-tuning, LP-FT and frame-level CNN and Transformer heads in the configurations we tried, plus a layer-wise analysis showing that dialect information peaks in the middle layers of both XLS-R and PhoWhisper.

II. Related Work

Vietnamese dialect identification. Earlier Vietnamese work mostly separated the three regional dialects. ViMD [1] introduced province-level labels for 63 provinces and benchmarked fine-tuned wav2vec 2.0, XLS-R, Whisper and PhoWhisper models. PhoWhisper-base was best with 0.411 macro-F1, and XLS-R-300m reached 0.373. The authors also found that geography influences which dialects are confused, which motivates our distance-aware objective.

Self-supervised speech representations. wav2vec 2.0 [4] and its multilingual version XLS-R [5] learn speech representa-

Table I
VIMD STATISTICS

Property	Value
Provinces (classes)	63 (pre-2025 boundaries)
Regions	North 25, Central 19, South 19
Utterances	18,949
train / valid / test	15,023 / 1,900 / 2,026
Audio	102.56 h, mean 19.5 s / utt.
Speakers	12,953 (88–309 per province)
Utterances per province	263–363 (median 303)
Chance accuracy	1.6%

tions from unlabelled audio. Whisper [6] is trained with weak supervision for transcription; PhoWhisper [7] fine-tunes it on 844 hours of Vietnamese speech, including in-house data from about 26,000 speakers across 63 provinces. The SUPERB benchmark [8] evaluates such models frozen, with a learned weighted sum of layers and a light downstream head; we follow this protocol. Layer-wise studies [9] show that different layers encode acoustic, phonetic and lexical information, so the best layer depends on the task.

Pooling and temporal architectures. For speaker and language recognition, attentive statistics pooling [10] and ECAPA-TDNN [11] are standard ways to summarise frame-level features. We adapt both as heads over frozen frames.

Fine-tuning versus probing. Kumar et al. [12] showed that fine-tuning a pretrained model with a randomly initialised head can distort its features, and proposed LP-FT: train the head first, then fine-tune. We test both strategies.

Distance-aware objectives. Label smoothing [13], [14] replaces one-hot targets with a mixture of the true label and a uniform distribution. Soft ordinal labels (SORD) [2] build targets from a distance between classes, hierarchy-aware losses [15] aim to make "better mistakes", and image geolocation uses "haversine smoothing", which decays the target of each geographic cell with its great-circle distance to the true location [3]. Geo-LS is a small variation on these ideas: it keeps label smoothing's fixed mass on the true class and distributes only the smoothed mass by geographic distance. We regard it as an adaptation rather than a new class of objective.

III. DATASET

We use ViMD [1] with its official splits. Table I summarises the statistics we measured on the released data.

What the label means. According to the dataset paper [1], the audio consists of interviews with local residents published by the 63 provincial television and broadcasting stations, mostly from news programmes; segments contain a single speaker, and news presenters were filtered out. Each recording is labelled with "the provincial dialect of the sample", which in practice is the province of the station. The paper does not describe a procedure that verified each speaker's origin, and it notes that some residents retain the dialect of a region where they previously lived. The label is therefore a noisy proxy for the speaker's dialect and is, by construction, tied to a broadcaster. ViMD also reports about three times as many male as female speakers. The data are released under CC BY-NC-ND 4.0.

Preprocessing. All audio was resampled to 16 kHz mono and stored as FLAC. The splits are intended to be speaker-disjoint [1], but we found two speakers in both the validation and test sets; we removed their validation utterances, leaving 1,898 validation utterances. The test set is untouched, so our numbers use the same test utterances as the original paper. Most speakers contribute only one or two utterances (about 1.5 on average).

Geography. Each province is represented by the coordinates of its capital city, and distances between provinces are great-circle (haversine) distances. The median distance from a province to its nearest neighbour is 44 km and the maximum is 113 km. Coordinates are approximate to about 5 km, which is small compared with these distances.

Transcripts. Each utterance comes with a transcript. We use transcripts only for text-only control baselines and to detect spoken place names, never as input to the main model. The most province-specific words in the transcripts are place names and local topics, consistent with the news-programme origin of the data.

IV. METHODOLOGY

Our system keeps pretrained speech encoders frozen and trains only a small head: (i) encode the waveform with a frozen encoder and keep the hidden states of every layer; (ii) mean-pool each layer over time, using only non-padded frames; (iii) standardise each layer with training-set statistics and combine the layers with learned softmax weights; (iv) classify the 63 provinces with a two-layer MLP (512 hidden units, GELU, dropout 0.3).

A. Frozen Backbones

XLS-R. We use facebook/wav2vec2-xls-r-300m [5] (24 Transformer layers, 1,024 dimensions). A validation-set probe showed that dialect information peaks around layer 5, so we keep the embedding output and layers 1 to 12 (13 states) and drop the rest. Long utterances are cut into 12-second windows; the pooled features of each window are averaged, weighted by frame count, over up to 60 seconds of audio.

PhoWhisper. We use the encoder of vinai/PhoWhisper-large [7] (32 layers, 1,280 dimensions), a Whisper model fine-tuned for Vietnamese ASR. Whisper always receives 30 seconds of padded log-mel input, so we pool only over frames that correspond to real audio. We keep all 33 states.

Fusion. For two backbones, each gets its own layer weights and standardisation; the two pooled vectors are concatenated (2,304 dimensions) before the MLP.

B. Layer-Weighted Pooling

For layer l with pooled vector $\mathbf{h}_l$, training mean $\boldsymbol{\mu}_l$ and standard deviation $\boldsymbol{\sigma}_l$, the utterance representation is

$$\mathbf{z} = \sum_{l=0}^{L} \frac{\exp(w_l)}{\sum_{k=0}^{L} \exp(w_k)} \cdot \frac{\mathbf{h}_l - \boldsymbol{\mu}_l}{\boldsymbol{\sigma}_l}. \tag{1}$$

The weights w are learned jointly with the MLP, as in SUPERB [8]. Because standardisation and mean-pooling are both linear, features can be pooled once and cached, so each head trains in seconds.

C. Geo-Aware Label Smoothing (Geo-LS)

Standard label smoothing gives the true province a mass of $1-\varepsilon$ and spreads ε uniformly over all provinces. Geo-LS keeps $1-\varepsilon$ on the true province but spreads ε over the other provinces in proportion to their geographic closeness. For true province i and distance $d(i,j)$ in kilometres,

$$q_i(j) = \begin{cases} 1-\varepsilon, & j = i, \\ \varepsilon \dfrac{\exp(-d(i,j)/\tau)}{\sum_{k \neq i} \exp(-d(i,k)/\tau)}, & j \neq i. \end{cases} \tag{2}$$

The off-diagonal terms sum to ε, so q_i is a valid distribution for any $\tau > 0$. The model is trained with cross-entropy against q_i. We select ε and τ on the validation set from $\varepsilon \in \{0.1, 0.2, 0.3\}$ and $\tau \in \{25, 50, 100, 200\}$ km; the chosen setting, $\varepsilon = 0.3$ and $\tau = 200$ km, lies at the upper edge of both grids, so larger values were not explored.

SORD [2] and haversine smoothing [3] derive the whole target, including the true class, from the distance kernel. Neighbouring provinces are only about 44 km apart, so a kernel wide enough to spread mass to neighbours also takes mass from the true class; with $\tau = 50$ km, less than half of the target mass remains on it. Geo-LS fixes the true-class mass independently of τ. The kernel uses straight-line distances between capitals, so it ignores natural dialect boundaries such as mountain passes, and provinces in dense regions spread their smoothing mass over more close neighbours than provinces at the edge of the country.

D. Frame-Level Architectures

To test whether temporal patterns such as tone contours add information beyond mean-pooled features, we extract frame-level features from XLS-R (layers combined with fixed weights taken from a linear probe, 1,024 dimensions, 25 frames per second, first 12 seconds) and train five heads on them: (a) mean pooling + MLP (control); (b) mean and standard deviation pooling + MLP; (c) attentive statistics pooling (ASP) [10]; (d) an ECAPA-style head [11] with three dilated Conv1D blocks with squeeze-excitation, followed by ASP (2.5 M parameters); (e) a two-layer Transformer encoder (4 heads, model dimension 256) followed by ASP (1.6 M parameters).

E. Fine-Tuning

We also fine-tune XLS-R end to end with label smoothing, keeping the convolutional feature encoder frozen and using SpecAugment-style time masking [16] ($p = 0.05$). Variant 1 fine-tunes all 24 layers with a randomly initialised head for 6 epochs on 6-second random crops (encoder learning rate 3×10^{-5}, head 10^{-3}). Variant 2 applies LP-FT [12] to the 12-layer encoder: the head is first trained on frozen features, then the whole model is fine-tuned for 12 epochs with a cosine schedule (encoder 2×10^{-5}, head 2×10^{-4}). Each variant was run once.

V. Experimental Setup

Training. Heads are trained with AdamW (learning rate 10^{-3}, weight decay 10^{-3}, batch 256) for up to 80 epochs, with early stopping on validation accuracy (patience 10). Every head configuration is run with seeds 0, 1 and 2. Hyperparameters and model choices are selected on the validation set; test results are reported for all configurations.

Baselines and controls. We train logistic-regression or linear-SVM models on (a) MFCC statistics (40 coefficients with deltas, mean and standard deviation, 240 dimensions); (b) the long-term average log-mel spectrum (LTAS, 80 dimensions); (c) TF-IDF word and character n-grams of the transcript; and (d) fine-tuned PhoBERT [17] on word-segmented transcripts. LTAS averages away most phonetic detail but still mixes recording channel, voice characteristics (including speaker sex) and the long-term distribution of speech sounds, so it is a partial proxy for the channel rather than a pure one. Controls (c) and (d) estimate how much spoken content alone reveals to these particular text models.

Metrics. We report accuracy and macro-F1 over the 63 provinces, top-5 accuracy, region accuracy (the predicted province lies in the true North/Central/South region), mean error distance in km (0 for a correct prediction) and the share of predictions within 100 km of the true province.

Single seeds versus ensembles. Two kinds of numbers appear in this paper and should not be mixed. Values written as mean $\pm$ s.d. are averages over the three seeds trained separately. Paired differences with confidence intervals, and all values in Tables V and VII, are computed for three-seed ensembles, whose predicted probabilities are averaged. Ensembling raises accuracy, so a difference between two ensembles need not equal the difference between the corresponding single-seed means; where both are relevant we report both.

Statistics. Confidence intervals use a paired bootstrap over test utterances with 1,000 to 2,000 resamples; we call a difference significant when its 95% interval excludes zero. These intervals reflect test-set sampling uncertainty for a fixed trained model or ensemble, not variation across training seeds, which is reported separately as the s.d. over seeds. Resampling single utterances treats them as independent, although some test speakers contribute several utterances and utterances from the same programme share a recording channel; the intervals may therefore be somewhat too narrow.

Reproducibility and compute. Code, configuration and all result files are available at https://github.com/stevedang2805/vimd-geo-dialect-id, including the province coordinates and the exact feature-extraction settings. All experiments ran on Google Colab Pro, mostly on an NVIDIA L4 (24 GB). Feature extraction took 15 to 75 minutes per backbone; each head trains in under a minute.

Table II
BASELINES ON THE TEST SET (SINGLE RUN; XLS-R USES ALL 25 STATES, THE FIRST 12 S AND CROSS-ENTROPY). ACCURACY WITH 95% BOOTSTRAP CI.

Model	Input	Acc. (%)	F1	Reg. (%)
LTAS + LR	Spectrum	19.7 [18.0, 21.5]	.182	56.1
MFCC stats + LR	Audio	24.6 [22.8, 26.5]	.235	65.4
TF-IDF + SVM	Text	21.9 [20.1, 23.6]	.208	58.1
PhoBERT (fine-tuned)	Text	17.8 [16.2, 19.4]	.133	67.1
XLS-R layer 4, linear	Audio	44.9 [42.7, 47.0]	.443	84.5
XLS-R wsum, linear	Audio	46.6 [44.4, 48.8]	.459	87.9
XLS-R wsum, MLP	Audio	**49.2** [47.2, 51.3]	**.486**	**89.4**

Table III
TEST ACCURACY (%) ON CONTROL SUBSETS (3-SEED ENSEMBLES; BASELINE: SINGLE RUN). "TEXT WRONG" AND "LTAS WRONG" ARE UTTERANCES MISCLASSIFIED BY THE TF-IDF AND LTAS BASELINES.

Subset	n	XLS-R base	PhoWhisper	Fusion D (LS)
All test	2,026	49.2	58.4	64.2
No place name	1,846	48.8	57.8	63.6
Own province named	151	50.3	62.3	67.6
Text wrong	1,583	45.9	55.5	61.3
LTAS wrong	1,627	44.0	52.1	58.1
Both wrong	1,290	41.0	49.2	55.4

VI. RESULTS

A. Baselines and Confound Controls

Frozen XLS-R features already reach 49.2% accuracy, more than twice any control baseline (Table II). The text-only and spectrum-only models are 11 to 16 times above chance, which confirms that ViMD contains province cues besides accent and that they must be controlled for.

Spoken content. Only 8.9% of test transcripts contain any province or capital-city name, and 7.5% contain the speaker's own province. Masking all place names lowers the TF-IDF model from 21.9% to 17.4%; the remainder comes from local topics and dialect words. Table III evaluates the audio models on subsets where these shortcuts are reduced. Accuracy without place names is within 0.7 points of the full test set for every model. On the 151 utterances that do name the own province, the text model jumps to 72.9%, while the audio models rise by only 1 to 4 points, within their wide confidence intervals. Both observations indicate that explicit place names are not a major source of the audio models' accuracy.

The "text wrong" subset should be read narrowly: it contains utterances whose transcript the TF-IDF model could not exploit, not utterances free of lexical information. A stronger text model, or the language knowledge inside PhoWhisper, might still extract cues from them.

Concern about PhoWhisper. PhoWhisper is trained for transcription, so it could exploit lexical content more than XLS-R. We find no evidence that it does: the share of accuracy it retains on the place-name-free, text-wrong and both-wrong subsets (98.9%, 94.9% and 84.1%) matches the XLS-R baseline (99.3%, 93.4% and 83.4%). Its ASR training data include in-house speech from 63 provinces whose source is not documented [7]; we cannot verify whether it overlaps with ViMD. Such data carry transcripts, not province labels, so possible overlap would expose the encoder to the same speech but would not leak the labels.

Recording channel. The LTAS baseline is 12 times above chance, so long-term spectra differ between provinces. Across provinces, LTAS accuracy and XLS-R accuracy are correlated (Spearman $\rho = 0.47$, $p < 0.001$ over 63 provinces); for example, Ninh Bình reaches 73% with LTAS alone. Channel cues therefore plausibly contribute part of the per-province accuracy. On the 1,290 utterances where both the text and the LTAS model fail, the fusion model keeps 55.4% with label smoothing (95% CI 52.7 to 58.1) and 54.4% with Geo-LS, about 35 times chance. Because LTAS is only a partial channel proxy, this shows that the accuracy does not depend on the cues LTAS captures, not that the model is free of channel information.

Table IV
TRAINING OBJECTIVES (XLS-R, 13 STATES, FIRST 12 S; MEAN ± S.D. OVER 3 SINGLE-SEED RUNS; PARAMETERS SELECTED ON VALIDATION)

Objective	Acc. (%)	F1	Reg. (%)	km	NLL
Cross-entropy	50.2 ± 0.5	.496	88.1	143	2.22
LS (ε=0.3)	**53.2** ± 0.6	**.526**	86.7	156	2.20
SORD (τ=5)	50.4 ± 0.3	.498	88.3	142	2.16
Geo-LS (0.3, 200)	51.3 ± 0.3	.507	**89.7**	**127**	**2.01**

B. Training Objectives

Label smoothing gives the highest accuracy, while Geo-LS gives the closest errors and the lowest negative log-likelihood (Table IV). These experiments use frozen XLS-R with 13 states and the first 12 seconds of audio. Comparing three-seed ensembles, Geo-LS beats cross-entropy on both accuracy ($+2.0$ points, 95% CI 0.6 to 3.5) and error distance (-16.3 km, CI -24.4 to -8.5). The corresponding difference between single-seed means in Table IV is smaller, $+1.1$ points (51.3% versus 50.2%).

SORD needs a tiny τ. With $\tau = 5$ km the SORD target is almost one-hot, since no neighbour lies that close ($e^{-44/5} < 10^{-3}$), and SORD behaves almost like cross-entropy ($+0.8$ points, CI -0.05 to 1.7). With $\tau = 50$ km on 25 layers it loses 7.7 points of accuracy against cross-entropy (CI -9.6 to -5.7) without reducing error distance; $\tau = 100$, 200 and 400 km are worse still (mean accuracy 35.0%, 27.8% and 24.3%). In this grid, no single kernel width kept the true class dominant while also spreading useful mass to neighbours; Geo-LS separates the two by construction.

Geo-LS versus label smoothing across settings. The geographic gain of Geo-LS is consistent, while its accuracy cost varies and is not significant on the strongest model (Table V).

A non-significant difference does not establish that two models are equivalent. For the fusion model, the 95% interval rules out accuracy losses larger than 1.5 points; we did not fix a non-inferiority margin in advance, so this is a post-hoc bound rather than a formal non-inferiority result.

Table V
GEO-LS MINUS LABEL SMOOTHING (PAIRED BOOTSTRAP ON 3-SEED ENSEMBLES, 95% CI)

Setting	Δ Acc. (points)	Δ Mean error (km)
XLS-R, 25 states	-1.2 $[-2.5, 0.2]$	-12.6 $[-22.1, -3.2]$
XLS-R, 13 states	-2.8 $[-4.1, -1.3]$	-14.4 $[-23.7, -5.2]$
XLS-R frame, mean head	-0.7 $[-2.0, 0.4]$	-18.7 $[-28.6, -9.5]$
Fusion (final)	-0.3 $[-1.5, 0.9]$	-8.2 $[-14.6, -1.4]$

Table VI
FINE-TUNING AND FRAME-LEVEL HEADS ON XLS-R WITH LABEL SMOOTHING. FROZEN ROWS: FRAME-LEVEL PIPELINE, MEAN ± S.D. OVER 3 SEEDS. FINE-TUNING ROWS: SINGLE RUNS.

Model	Params	Acc. (%)	km	Reg. (%)
Frozen, mean + MLP (ref.)	0.56 M	**52.7** ± 1.1	156	86.4
Frozen, ECAPA-style + ASP	2.5 M	48.0 ± 0.8	141	88.5
Frozen, Transformer + ASP	1.6 M	47.3 ± 0.5	**133**	**89.6**
Full FT, 24 layers, 6 ep.	∼300 M	40.7	244	79.0
LP-FT, 12 layers, 12 ep.	∼150 M	no gain over LP (see text)		

Fewer layers are better. Keeping only the embedding output and layers 1 to 12 of XLS-R improves the single-seed mean of every objective over using all 25 states: cross-entropy from 49.7% to 50.2%, label smoothing from 51.6% to 53.2%, Geo-LS from 50.4% to 51.3%. The upper layers add noise for this task; Section VII-A shows why.

C. Fine-Tuning and Frame-Level Architectures

In our configurations, neither fine-tuning nor a more complex head beat a frozen encoder with mean pooling and an MLP (Table VI). The frozen rows of Table VI come from the frame-level pipeline (Section IV-D), which uses fixed layer weights; its mean-pooling reference (52.7%) is therefore not identical to the learned-weight model of Table IV (53.2% with label smoothing).

Full fine-tuning. With a random head, validation accuracy started at 15% after the first epoch and was still rising after epoch 6 (43.5%), so this run was under-trained. Its best test accuracy (40.7%) stayed far below the frozen baseline. A longer schedule might narrow the gap; we did not test it.

LP-FT. Starting from a trained head, the model at epoch 0 reproduces the frozen result (validation 55.7%), which confirms that the pipelines match. One epoch of fine-tuning then drops validation accuracy to 46.0% (37.2% with Geo-LS), and it recovers only to 50.9% after 11 epochs, never reaching the starting point. Training loss kept falling, which suggests that with about 240 utterances per province the encoder adapts to speaker and recording specifics rather than to accent. Parameter-efficient fine-tuning and other learning-rate schedules remain untested.

Temporal heads. Comparing three-seed ensembles, the Conv1D and Transformer heads lose 3.7 and 4.3 points of accuracy against mean pooling (CIs -5.5 to -1.8 and -6.3 to -2.5) but bring errors 17.9 and 26.5 km closer; the corresponding differences between single-seed means in Table VI are 4.7 and 5.4 points and 15 and 23 km. They also raise region accuracy by 2 to 3 points. They capture coarse regional patterns but not the finer distinctions between neighbouring provinces. In an earlier run with 512-dimensional PCA features and 6-second crops, adding standard-deviation statistics hurt further (statistics pooling 42.5%, ASP 38.6%, against 49.0% for mean pooling), possibly because frame-level variance encodes speaker and channel rather than dialect.

Takeaway. With this amount of data, changing the training target brought the geographic benefit of a temporal head at a third of its parameters and without its accuracy cost: in the same frame-level pipeline, mean pooling with Geo-LS reaches 52.2% and 128 km (single-seed means), against 47.3% and 133 km for the Transformer.

D. Backbone Fusion and Final Model

Fusing XLS-R and PhoWhisper is the largest single improvement in this study: $+6.7$ points over XLS-R alone (CI 5.1 to 8.3) and $+5.7$ over PhoWhisper alone (CI 4.1 to 7.5). Using the whole utterance instead of the first 12 seconds helps XLS-R only slightly ($+1.1$ points, CI -0.4 to 2.6). Table VII reports three-seed ensembles; the last row averages all nine heads of B, C and D. Single-seed means are about 3 points lower (fusion with LS: $61.1 \pm 1.2\%$; with Geo-LS: $61.0 \pm 0.6\%$).

We take the fusion architecture with Geo-LS as the main system: it is within 0.7 points of the nine-head ensemble, whose gain over the fusion ensemble is not significant ($+0.4$ points, CI -0.6 to 1.5), and it makes the geographically closest errors. Against our first frozen XLS-R baseline (49.2%, a single run), it improves accuracy by 14.7 points (CI 12.6 to 16.8) and reduces the mean error by 60 km (CI 50 to 71).

Comparison with the ViMD paper. Table VIII lists macro-F1 together with the fine-tuned baselines published with the dataset. The comparison is not controlled. The paper fine-tunes each model and splits recordings longer than 10 seconds into segments of at most 10 seconds; it does not state how segment predictions are aggregated, so its evaluation unit may differ from our whole utterances. Its best model is PhoWhisper-base (74 M parameters), whereas our systems use frozen XLS-R-300m and the PhoWhisper-large encoder, and their hyperparameters were tuned on our pipeline. The difference therefore combines backbone size, input length, evaluation unit and the frozen-probing protocol, and cannot be attributed to any one of them; in particular, Geo-LS does not improve macro-F1 over label smoothing.

VII. ANALYSIS

A. Where Dialect Lives in the Encoders

Dialect information peaks in the middle of both encoders and fades towards the output (Fig. 1). For XLS-R, a linear probe on single layers is best at layer 5 (51.0% validation accuracy). For PhoWhisper, it is best at layer 16 of 32 (53.4%) and falls to 38.0% at the last layer, which Whisper shapes for transcription. This explains why dropping the upper half of XLS-R helps (Section VI-B). The learned layer weights agree: XLS-R puts most of its weight on layers 3 to 6, and the PhoWhisper-only model on layers 10 to 20. In the fusion

Table VII
FINAL MODELS ON THE VIMD TEST SET (3-SEED ENSEMBLES; LAST ROW: ALL 9 HEADS OF B, C AND D; 95% BOOTSTRAP CI)

Model	Accuracy (%)	Macro-F1	Top-5 (%)	Region acc. (%)	Mean error (km)	Within 100 km (%)
B: XLS-R, full utterance, LS	57.5 [55.2, 59.6]	0.569	81.1	88.9	127	69.2
C: PhoWhisper-large, LS	58.4 [56.2, 60.5]	0.575	84.3	93.0	98	71.9
D: Fusion B + C, LS	64.2 [62.0, 66.2]	0.635	85.6	93.3	88	75.8
D: Fusion B + C, Geo-LS	63.9 [61.6, 65.9]	0.631 [0.608, 0.649]	86.3	**94.4**	**80**	76.4
Ensemble of B, C and D, LS	**64.6** [62.4, 66.6]	**0.637** [0.615, 0.654]	**86.5**	93.8	82	**76.6**

Table VIII
MACRO-F1 ON THE VIMD 63-PROVINCE TEST SET. VIMD ROWS USE A DIFFERENT PROTOCOL (SEE TEXT).

System	Approach	Macro-F1
wav2vec2-large-xlsr-53 [1]	Fine-tuned, $\leq$10 s seg.	0.326
wav2vec2-xls-r-300m [1]	Fine-tuned, $\leq$10 s seg.	0.373
PhoWhisper-base [1]	Fine-tuned, $\leq$10 s seg.	0.411
Ours: XLS-R, 13 states, LS	Frozen; seed mean	0.527
Ours: fusion, Geo-LS	Frozen; 3-seed ens.	0.631
Ours: ensemble of B, C and D	Frozen; 9 heads	**0.637**

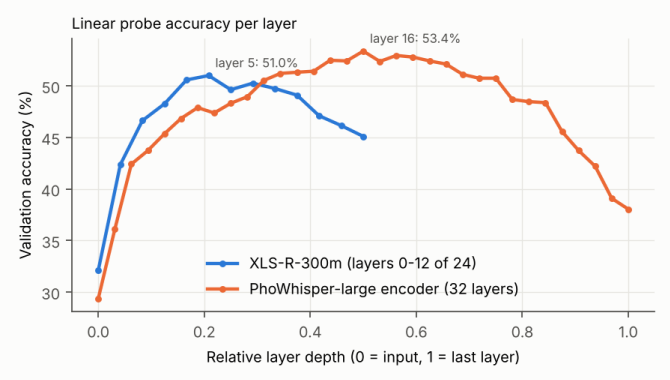


Figure 1. Linear-probe validation accuracy for each layer, plotted against relative depth. Only layers 0 to 12 of XLS-R were extracted.

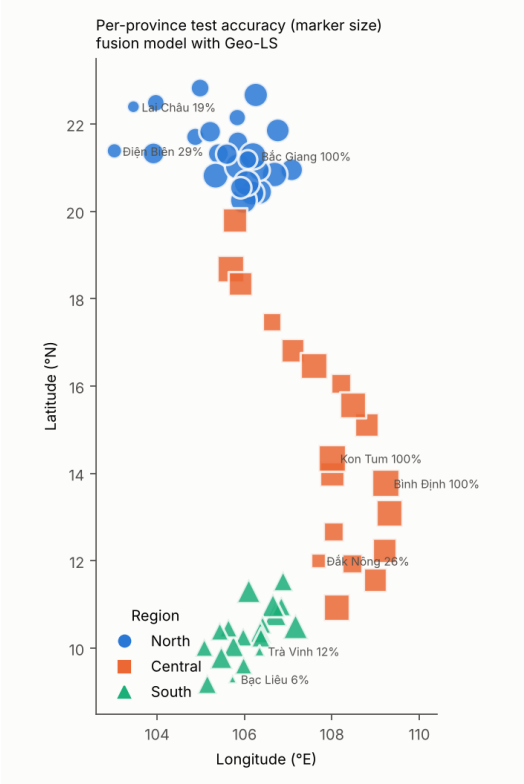


Figure 2. Per-province test accuracy of the fusion model with Geo-LS. Each marker sits at the province capital; its size is proportional to accuracy.

model, PhoWhisper also gives noticeable weight to its last two layers, possibly because they carry information that XLS-R lacks.

B. Errors Follow Geography

When the model is wrong, it is usually wrong nearby. For the first XLS-R baseline, trained with plain cross-entropy, wrong predictions were on average 267 km from the true province, against 682 km for a random wrong guess, and the number of confusions between two provinces falls with their distance (Spearman $\rho = -0.34$ over ordered province pairs). For the final Geo-LS model the corresponding value is $\rho = -0.33$, and 76% of all test utterances are placed within 100 km of the true province. Province pairs that share a province are not independent, so ordinary Spearman p-values are not valid here; significance should be assessed with a Mantel permutation test [18], which permutes the province labels of the distance matrix. Geo-LS trains the model to prefer nearby provinces, so evidence about the data must come from the cross-entropy model. Its pattern is consistent with gradual geographic variation in accent, but neighbouring provinces may also share broadcasting styles or recording conditions, which would produce the same pattern.

C. Easy and Hard Provinces

Per-province accuracy of the fusion model with Geo-LS ranges from 6% to 100% (Fig. 2). The hardest provinces are Bạc Liêu (6%), Trà Vinh (13%), Lai Châu (19%), Đắk Nông (26%) and Điện Biên (29%); Bắc Giang, Bình Định and Kon Tum reach 100%. Each province has only about 32 test utterances, so these values are imprecise. A plausible reading is that the hard provinces sound very close to their neighbours in the Mekong Delta, or have mixed migrant populations in the highlands; we have not verified this. Some of the easiest provinces, such as Ninh Bình, also have a distinctive recording channel, so their scores may be partly inflated.

VIII. Limitations

Label validity and recording source. Labels denote the province of the broadcasting station, not a verified speaker origin, and each province's audio comes from one broadcaster. Our controls show that explicit place names and the cues captured by text and LTAS baselines explain little of the accuracy, but they cannot separate accent from broadcaster-specific recording conditions. Results should be read as accuracy on ViMD-like broadcast recordings, not as speaker-origin identification from arbitrary audio.

Speaker sex. ViMD has about three times as many male as female speakers. A binary attribute cannot by itself identify many provinces, but we have not tested whether accuracy differs between male and female speakers or whether the sex ratio varies across provinces.

No out-of-domain test. All results come from the ViMD test split. We have not evaluated on independently collected recordings, such as phone speech from speakers of known origin, where we expect a substantial drop.

Small data per class. With about 240 training utterances per province, fine-tuning and larger heads overfit in our runs. The negative fine-tuning results hold for the two schedules we tried, each run once, not for fine-tuning in general; parameter-efficient methods such as LoRA or adapters remain untested.

Statistics. Bootstrap intervals resample utterances and reflect only test-set uncertainty; seed variability is reported separately. Utterances that share a speaker or programme are treated as independent, so intervals may be slightly too narrow. Our calibration evidence is limited to negative log-likelihood; expected calibration error and Brier score were not measured.

Approximate geography. Each province is reduced to the coordinates of its capital, and labels follow the 63-province boundaries in force before the 2025 merger. The Geo-LS kernel ignores natural dialect boundaries, and its chosen parameters lie at the edge of the search grid.

Validation–test gap. Validation accuracy is about 6 points above test accuracy for every model, so comparisons between models are consistent, but absolute validation numbers are optimistic. We draw conclusions from test results only.

Comparison with prior work. The ViMD baselines use a different protocol (fine-tuning, smaller backbones, segments of at most 10 seconds), so the gain over them combines several factors.

IX. Conclusion

Frozen self-supervised encoders with a learned layer-weighted head identify the broadcasting province of Vietnamese speech in ViMD with 63.9% accuracy for a three-seed ensemble of a fused XLS-R and PhoWhisper head (61.0% for individual seeds), and 64.6% for a nine-head ensemble. This is well above the 0.411 macro-F1 published with the dataset, although the protocols differ. Controls with transcripts, long-term spectra and place-name-free subsets indicate that the accuracy is not explained by explicit place names or by the content and spectral cues our baselines capture; because labels are tied to broadcasters, residual recording-source effects cannot be excluded. Geo-aware Label Smoothing reduced the geographic error of every model we tried by 8 to 19 km against standard label smoothing; on the strongest model the accompanying accuracy change was -0.3 points (95% CI -1.5 to 0.9). In this low-resource regime and in our configurations, the frozen-feature approach outperformed full fine-tuning, LP-FT and temporal CNN and Transformer heads.

Four directions follow. First, testing on independently recorded speech from speakers of known origin would measure how far the system generalises beyond ViMD's broadcast conditions. Second, speaker-level resampling, permutation tests for the geographic pattern, calibration metrics and a breakdown by speaker sex would tighten the statistical evidence. Third, adversarial training against speaker or channel information could reduce the remaining confound. Fourth, parameter-efficient fine-tuning, and joint dialect identification and speech recognition using ViMD's transcripts, could show whether adaptation helps once overfitting is controlled.

References

[1] N. Van Dinh, T. C. Dang, L. T. Nguyen, and K. Van Nguyen, "Multi-dialect Vietnamese: Task, dataset, baseline models and challenges," in *Proc. EMNLP*, 2024, arXiv:2410.03458.

[2] R. Díaz and A. Marathe, "Soft labels for ordinal regression," in *Proc. IEEE CVPR*, 2019.

[3] L. Haas, M. Skreta, S. Alberti, and C. Finn, "PIGEON: Predicting image geolocations," in *Proc. IEEE/CVF CVPR*, 2024.

[4] A. Baevski, H. Zhou, A. Mohamed, and M. Auli, "wav2vec 2.0: A framework for self-supervised learning of speech representations," in *Proc. NeurIPS*, 2020.

[5] A. Babu, C. Wang, A. Tjandra *et al.*, "XLS-R: Self-supervised cross-lingual speech representation learning at scale," in *Proc. Interspeech*, 2022.

[6] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, "Robust speech recognition via large-scale weak supervision," in *Proc. ICML*, 2023.

[7] T.-T. Le, L. T. Nguyen, and D. Q. Nguyen, "PhoWhisper: Automatic speech recognition for Vietnamese," in *Proc. ICLR Tiny Papers*, 2024, arXiv:2406.02555.

[8] S. Yang, P.-H. Chi, Y.-S. Chuang *et al.*, "SUPERB: Speech processing universal performance benchmark," in *Proc. Interspeech*, 2021.

[9] A. Pasad, J.-C. Chou, and K. Livescu, "Layer-wise analysis of a self-supervised speech representation model," in *Proc. IEEE ASRU*, 2021.

[10] K. Okabe, T. Koshinaka, and K. Shinoda, "Attentive statistics pooling for deep speaker embedding," in *Proc. Interspeech*, 2018.

[11] B. Desplanques, J. Thienpondt, and K. Demuynck, "ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification," in *Proc. Interspeech*, 2020.

[12] A. Kumar, A. Raghunathan, R. Jones, T. Ma, and P. Liang, "Fine-tuning can distort pretrained features and underperform out-of-distribution," in *Proc. ICLR*, 2022.

[13] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna, "Rethinking the Inception architecture for computer vision," in *Proc. IEEE CVPR*, 2016.

[14] R. Müller, S. Kornblith, and G. Hinton, "When does label smoothing help?" in *Proc. NeurIPS*, 2019.

[15] L. Bertinetto, R. Mueller, K. Tertikas, S. Samangooei, and N. A. Lord, "Making better mistakes: Leveraging class hierarchies with deep networks," in *Proc. IEEE CVPR*, 2020.

[16] D. S. Park, W. Chan, Y. Zhang *et al.*, "SpecAugment: A simple data augmentation method for automatic speech recognition," in *Proc. Interspeech*, 2019.

[17] D. Q. Nguyen and A. T. Nguyen, "PhoBERT: Pre-trained language models for Vietnamese," in *Findings of EMNLP*, 2020.

[18] N. Mantel, "The detection of disease clustering and a generalized regression approach," *Cancer Research*, vol. 27, no. 2, pp. 209–220, 1967.